

NYC Property Sales - Harvard EdX Data Science Capstone

Mary Lampmann

12/29/2020

Introduction

The City of New York has made a dataset of all Property Sales transactions in the city for the twelve month period from September 2016 to September 2017 available for data analysis (Source:Kaggle.com).

#Kaggle webpage * [NYC Property Sales]<https://www.kaggle.com/new-york-city/nyc-property-sales>

#Github link for download of Kaggle csv file <https://raw.githubusercontent.com/MLamp20/NYCProps/main/nyc-rolling-sales.csv>

This report documents the application of machine learning (ML) techniques to the variables in that dataset in order to predict the Sale Price of each property as accurately as possible. This report will document my analysis and present my findings with supporting statistics and figures.

The goal of this project is the training of multiple ML algorithms which will use inputs from a subset of the data to predict Sale Price in a separate subset of that same data, determining the algorithm that does so with the lowest residual mean squared error (RMSE). I will be standardizing the data, both to eliminate skew in the Sale Price and to modify the dataset for use with multiple types of machine learning algorithms, so the target number for the RMSE will be the lowest RMSE below the standard deviation of 1.0.

Project Key steps

1. Preprocess the dataset, inclusive of:
 - a. Renaming or creation of variables
 - b. Class conversions: Character to Numeric, Character to Time
 - c. Removal or replacement of blanks, NA, and zeroes as appropriate
 - d. Removal of duplicated lines
2. Conduct exploratory data analysis (EDA) on the training set **nyc** for use in machine learning (ML) model development, as well as additional dataset modification based on EDA learnings, inclusive of:
 - a. Log transform and scale variables to remove skew identified in EDA
 - b. Remove extreme outliers from dataset
 - c. Conversion of variables to factors
 - d. One-hot conversion of factors and remaining character classes
 - e. Eliminate variables with lowest correlation to Sale Price to streamline model processing times
3. After EDA, partition the **nyc** training dataset to allow separate training (*train*) and test sets (*test*) for use in evaluation of ML models
4. Iteratively train different machine learning algorithms on the dataset to elicit the lowest RMSE results, identifying the final recommended process

Methods/Analysis

Exploratory Data Analysis New York City itself is a big place. The dataset, **nyc**, is comprised of **84548** rows and **22** columns, and includes distinct records for **5** boroughs (1-Manhattan, 2-Bronx, 3-Brooklyn, 4-Queens, 5-Staten Island), **254** neighborhoods, **11566** blocks, **47** building class categories, and at the point of sale, **166** building classes. Additional variables in the dataset include the address, apartment number (where applicable), zip code, parcel lot number, residential, commercial and total units in the property, square footage (gross and for the land parcel lot), year built, tax class of the property, sale price, and date of sale.

Preprocessing The first step in preprocessing is simplification of 16 of the 22 variable names for ease of reading, both in the removal of spaces and abbreviation of certain terms.

Evaluation of the distinct values for each variable is the next step, which immediately identifies a column **EASEMENT** with no values. We will extract this from the dataset in a future step.

	ADDRESS	X1	BLOCK	SALE_PX	LAND_SF	GRSS_SF	APT_NO	LOT	SALE_DATE
1	67563	26736	11566	10008	6062	5691	3988	2627	364
	NEIGHBORHOOD	TTL_UNITS	ZIP	RESID_UNITS	BLDG_CL_0917	BLDG_CL_SOLD	BUILT		
1	254	192	186	176	166	166	158		
	COMM_UNITS	BLDG_CL_CATEGORY	TAX_CL_0917	BOROUGH	TAX_CL_SOLD	EASEMENT			
1	55		47	10	5	4	0		

Next up: a review of the percentage of zero values in each vector. Biggest issues exist in the UNITS variables (a notable 94% of the commercial units (COMM_UNITS)), the SF variables (gross square foot (GRSS_SF) and land square foot (LAND_SF)), and in Sale Price (SALE_PX). We will have to remove records with zero Sale Price, but will look to replace other zero values with mean values in future steps.

COMM_UNITS	RESID_UNITS	TTL_UNITS	GRSS_SF	LAND_SF	SALE_PX
94	29	23	14	12	12
BUILT	ZIP				
8	1				

Potentially more impactful to our dataset than zeroes are the percentage of NA values. Surprisingly, a review of the NA values identifies only 4 vectors with NA values. We see that 77% of Apartment Number vector (**APT_NO**) values are NAs, suggesting minimal value in using that variable in our model building. We will look to replace NAs in “at present”(0917) values of the tax class and building classes of the properties with the values of those classes at the point of sale.


EASEMENT	APT_NO	TAX_CL_0917	BLDG_CL_0917
100	77	1	1

The original dataset consisted of character, numeric, and logical vectors. Next step in our preprocessing is conversion of character vectors that represent numbers to numeric vectors, or, in the case of date of sale (SALE_DATE), to a time vector.

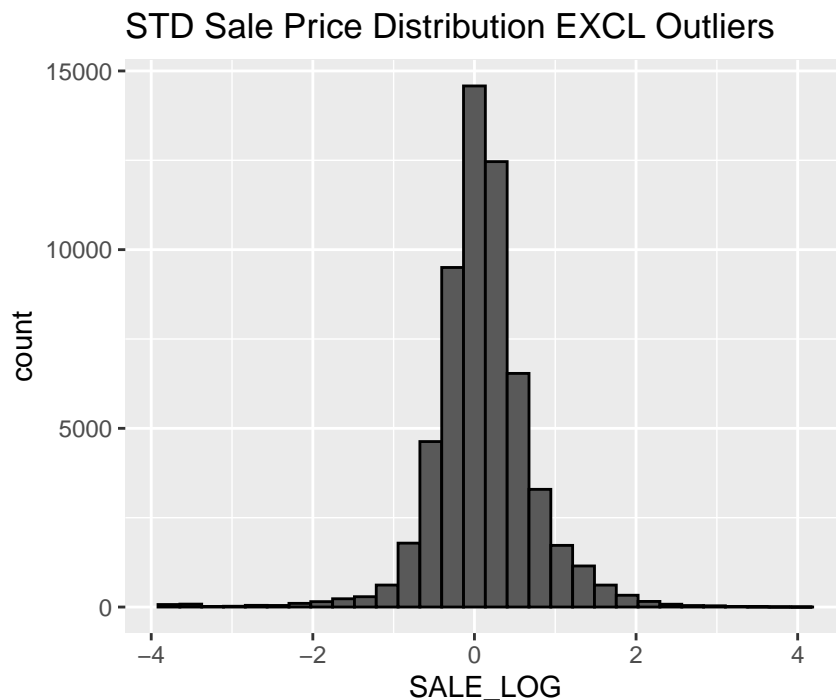
Once we have completed that step, we conduct our replacements of zero and NA values, and remove records where the sale price is NA, blank, or zero. We have also identified duplicated lines in the dataset, so will remove those, the EASEMENT column noted above, and then begin some data visualization.

With completion of the above, our revised dataset is now **59592 rows** and **21 columns** in dimension.

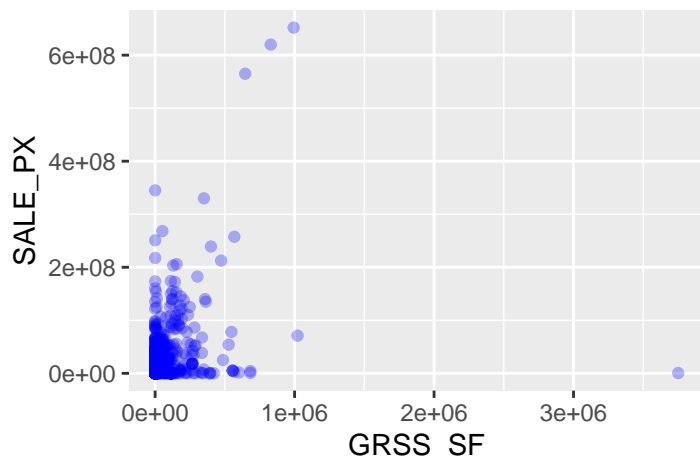
Data Visualization We start with a review of the distribution of our Sale Prices (SALE_PX). We have a pretty wide range of Sale Prices, starting at \$1 and maxing out at \$2.21 Billion, so we will first review this via a log 10 transformation of the plot.



The data is clearly skewed with a long tail. Next step: log transform and then standardize the sale price. Once complete, I elected to remove the records that have an absolute scaled Sale Log value of >4 , which leaves us with the following plot and the variable SALE LOG, which represents the log transformed and scaled Sale Prices:




Square Footage is a metric that is often discussed in the context of sale prices, here's the comparison of the gross square footage (GRSS_SF) of our data set with the Sale Price:



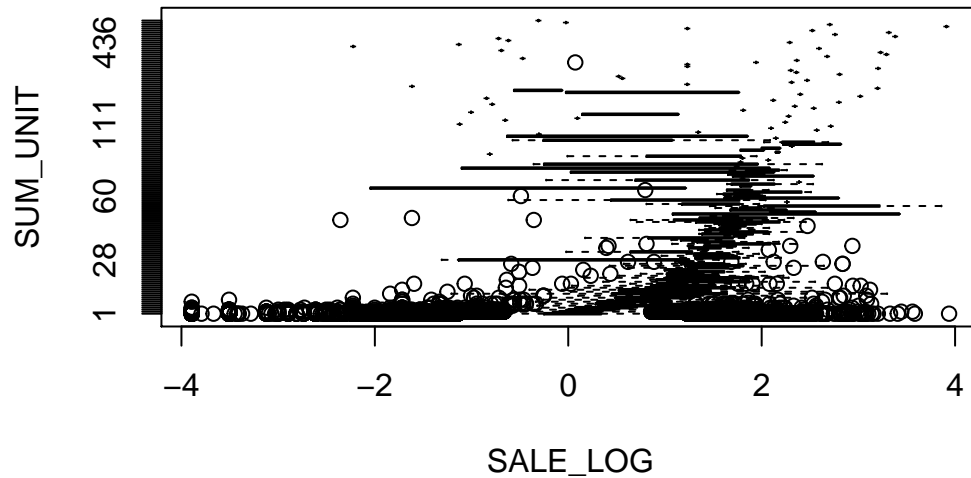
The values are tightly clustered, and we will remove that GRSS_SF extreme outlier on the far right in a future step.

Sales Price Boxplot Evaluating sales prices by category, we see Condo Warehouses/Factory with the lowest average Sales Price, and Luxury Hotels with prices at the top of the group.

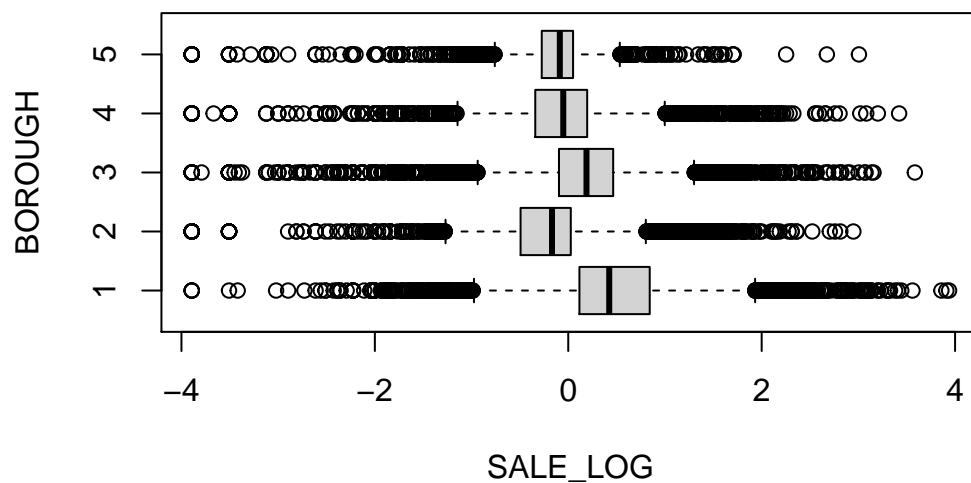


Although there appears to be some linearity in the 18-50 SUM UNIT counts, the sale price averages for the


sum of units are notably variable in the 50-100 unit range.



In sale price by Borough, we see that Borough 1, Manhattan has both the highest average price and the largest interquartile range (IQR). Borough 2, the Bronx has the lowest average, and Borough 5, Staten Island has the smallest interquartile range.



Time Effect on Sale Price by Borough The 5 boroughs were faceted for analysis as to whether time played a significant role in Sale Price. Although there was some time effect in Borough #1, Manhattan, most notable in the October-January time period, the other 4 boroughs average sale prices were remarkably stable across the weeks of the dataset. Considering this, time will not be incorporated into the ML model.



We've reviewed visualizations on sale price vs. some of the variables, but now let's evaluate the count of sales transactions (n), neighborhoods ($nbrhd$), blocks (blk), building class categories ($categ$), and average sale price per gross square footage (avg_pxgsf). Note that avg_px_K is an expression of the Sale Price in Thousands.

	BOROUGH	n	nbrhd	categ	blk	avg_px_K	avg_pxgsf
1	1	14243	39	39	1286	3147	1332
2	2	5037	38	34	1732	826	214
3	3	15349	60	43	4714	1307	607
4	4	18115	59	39	6590	751	395
5	5	5868	58	27	2238	555	273

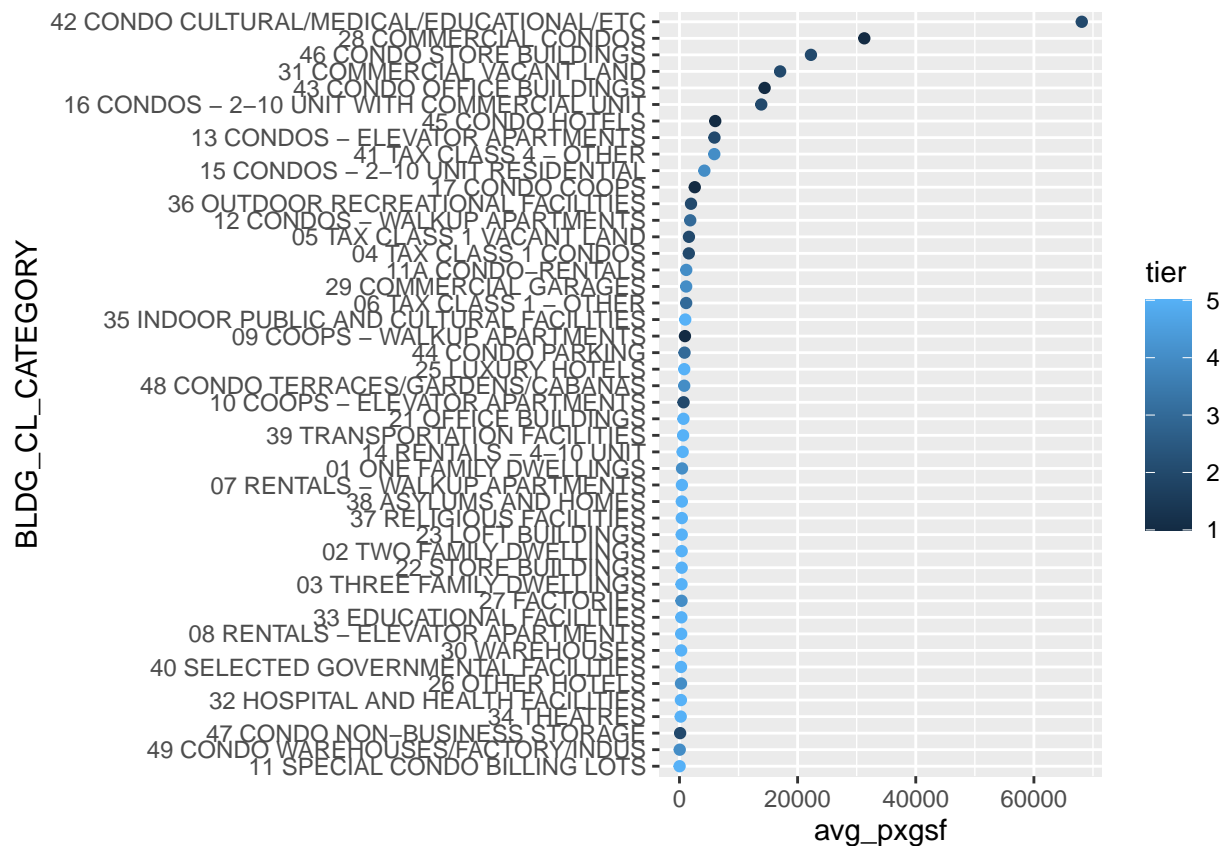
The greatest number of sales transactions came from Boroughs 4,3, and 1. The number of transactions in Borough 4 and 3 are the largest, but both Boroughs rank at the top in number of blocks, and theoretically, have the largest inventory of properties. Borough #1, Manhattan, had the highest average price per gross square footage of property sold at \$1332.

Building Class Categories Slicing the information by building class category, we see that the top 10 categories by number of transactions (n) are comprised of family dwellings, apartments, and condos. These categories represent 92% of the transactions in our entire dataset. The categories in lines 1-8 are present in each of the 5 boroughs.

	BLDG_CL_CATEGORY	n	boro	nrhd	blk	avg_px_K	avg_pxgsf
1	01 ONE FAMILY DWELLINGS	12701	5	219	6317	670	434
2	10 COOPS - ELEVATOR APARTMENTS	11513	5	141	1426	801	690
3	13 CONDOS - ELEVATOR APARTMENTS	10262	5	129	1117	2100	5922
4	02 TWO FAMILY DWELLINGS	9883	5	215	5474	802	375
5	09 COOPS - WALKUP APARTMENTS	2503	5	114	699	494	931
6	03 THREE FAMILY DWELLINGS	2322	5	164	1715	1008	349
7	07 RENTALS - WALKUP APARTMENTS	1742	5	171	1290	3444	413
8	04 TAX CLASS 1 CONDOS	1248	5	111	401	561	1582
9	17 CONDO COOPS	1111	4	47	175	995	2592
10	15 CONDOS - 2-10 UNIT RESIDENTIAL	1033	3	72	431	1495	4216

We do see some notable swings in average sale price per square foot here, but in further analysis, note that most of the condo categories had large percentages of NAs for the gross square footage that we use in this calculation, and it is likely that in using the average SF for all categories, our assigned values are too low for the condo categories specifically.

The categories have been assigned tiers corresponding to the percentage of GRSS_SF records that were NA, with tier 1 representing 75+% NA, tier 2 50-74%, tier 3 25-49%, tier 4 1-24%, and tier 5 with no NA in GRSS_SF.



If we rework our sort and examine the highest average price per gross SF (avg_pxgsf), we clearly see that those same condos where we have replaced 50%+ of GRSS_SF have extremely high avg_pxgsf averages. In a future iteration of this analysis, we may want to replace the average square footage assignments we have made for NAs to the average for each specific building class category.

	BLDG_CL_CATEGORY	n	avg_pxgsf	na_pct	tier
1	42 CONDO CULTURAL/MEDICAL/EDUCATIONAL/ETC	4	68108	0.60	2
2	28 COMMERCIAL CONDOS	14	31292	0.79	1
3	46 CONDO STORE BUILDINGS	74	22256	0.70	2
4	31 COMMERCIAL VACANT LAND	187	17039	0.54	2
5	43 CONDO OFFICE BUILDINGS	241	14419	0.87	1
6	16 CONDOS - 2-10 UNIT WITH COMMERCIAL UNIT	38	13847	0.61	2
7	45 CONDO HOTELS	76	6058	0.98	1
8	13 CONDOS - ELEVATOR APARTMENTS	10262	5922	0.74	2
9	41 TAX CLASS 4 - OTHER	55	5893	0.23	4
10	15 CONDOS - 2-10 UNIT RESIDENTIAL	1033	4216	0.25	4

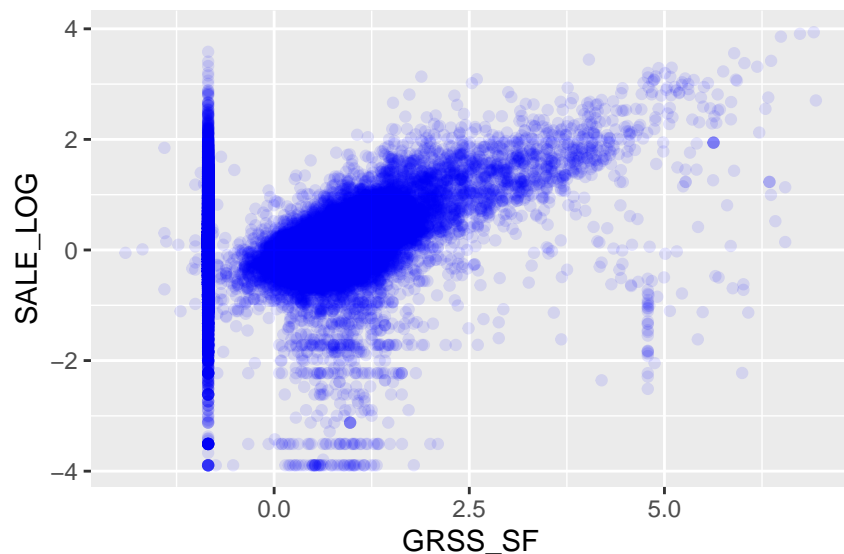
Resorting the categories by total US dollars sold, the top 10 categories here represent almost 81% of the total dollar sales in our dataset. The distortion on the average price per square foot as a result of the assigned GRSS_SF value for condo categories continues to be notable on this chart.

	BLDG_CL_CATEGORY	TTL_SLS_K	avg_pxgsf
1	13 CONDOS - ELEVATOR APARTMENTS	21552283	5923
2	10 COOPS - ELEVATOR APARTMENTS	9226653	690
3	01 ONE FAMILY DWELLINGS	8505244	434
4	02 TWO FAMILY DWELLINGS	7925747	375
5	07 RENTALS - WALKUP APARTMENTS	5999432	413
6	21 OFFICE BUILDINGS	5464142	672
7	08 RENTALS - ELEVATOR APARTMENTS	4910787	294
8	03 THREE FAMILY DWELLINGS	2340114	349
9	22 STORE BUILDINGS	1975404	374
10	15 CONDOS - 2-10 UNIT RESIDENTIAL	1544434	4216

Categories that make up the bottom of the list generally appear to be non “living” categories - warehouses, transportation hubs, parking, and storage.

	BLDG_CL_CATEGORY	TTL_SLS_K	avg_pxgsf
1	49 CONDO WAREHOUSES/FACTORY/INDUS	164	46
2	39 TRANSPORTATION FACILITIES	220	620
3	11 SPECIAL CONDO BILLING LOTS	600	20
4	47 CONDO NON-BUSINESS STORAGE	2730	110
5	48 CONDO TERRACES/GARDENS/CABANAS	3414	802
6	40 SELECTED GOVERNMENTAL FACILITIES	16000	261
7	06 TAX CLASS 1 - OTHER	25116	1135
8	36 OUTDOOR RECREATIONAL FACILITIES	56946	1952
9	44 CONDO PARKING	83483	856
10	42 CONDO CULTURAL/MEDICAL/EDUCATIONAL/ETC	96603	68107

Given our learnings thus far, we are ready to move on to log transformation and standardization of the square footage and units vectors. One more outlier trim on gross square footage, and here is our revised plot of that vs. Sale Price, demonstrating a linear relationship:



Last steps before moving onto partitioning will be conversion of Borough and Tax Class Sold to factors, and then One Hot Encoding of both vectors as well as all the remaining character vectors we want to consider in our model.

Here are the final list of variables that we will use in our model, and their correlation with the Sale Price (logged and scaled, SALE_LOG). Note that GRSS_SF with 19.3% correlation to SALE LOG is not the highest correlated variable. The highest correlation instead is whether the property is in Borough.1, Manhattan.

	term	SALE_LOG
1	BOROUGH.1	0.366
2	BOROUGH.2	-0.162
3	BOROUGH.4	-0.201
4	BOROUGH.5	-0.135
5	NEIGHBORHOODCIVIC.CENTER	0.126
6	NEIGHBORHOODTRIBECA	0.120
7	BLDG_CL_CATEGORY01.ONE.FAMILY.DWELLINGS	-0.118
8	BLDG_CL_CATEGORY07.RENTALS...WALKUP.APARTMENTS	0.154
9	BLDG_CL_CATEGORY08.RENTALS...ELEVATOR.APARTMENTS	0.161
10	BLDG_CL_CATEGORY09.COOPS...WALKUP.APARTMENTS	-0.132
11	BLDG_CL_CATEGORY10.COOPS...ELEVATOR.APARTMENTS	-0.153
12	BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS	0.258
13	BLDG_CL_CATEGORY44.CONDO.PARKING	-0.169
14	BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS.STORAGE	-0.110
15	GRSS_SF	0.193
16	TAX_CL_SOLD.1	-0.143
17	TAX_CL_SOLD.2	0.113
18	BLDG_CL_SOLD1	0.123
19	BLDG_CL_SOLD6	-0.132
20	BLDG_CL_SOLD7	0.105
21	BLDG_CL_SOLD1	0.106
22	BLDG_CL_SOLD4	-0.161
23	BLDG_CL_SOLDRG	-0.149
24	SUM_UNIT	0.234

Model Construction Methods

Split nyc into train and test sets The next step in the machine learning algorithm creation was the split of the **nyc** data set into training and testing subsets, *train* and *test* respectively.

I chose to use an 80/20 split on the partitioning based on the Pareto principle, and because I wanted to leave a testing set of sufficient size for the successful testing of a variety of machine learning algorithms.

I used R's RMSE function to calculate the RMSE, as well as a data frame that would detail the RMSE values associated with different machine learning algorithms.

```
#Test set will be 20% of data as set up for multiple algorithm uses later
set.seed(1,sample.kind = "Rounding")
test_index<-createDataPartition(y=nyc_pre_split$SALE_LOG,times=1, p=0.2, list=FALSE)
train<-nyc_pre_split[-test_index,]
test<-nyc_pre_split[test_index,]
```

The model fittings machine learning algorithms trained: Next steps were the selection of different machine learning algorithms for model fittings.

Linear Model The plot we viewed earlier with gross square footage (GRSS_SF) and Sale Price (SALE_LOG) indicated a strong linear relationship, so the logical place to start our machine learning exercise is with simple linear regression.

```
model<-lm(SALE_LOG ~ ., data = train)
print(tidy(model),n=Inf)
```

A tibble: 25 x 5

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	0.386	0.0126	30.6	2.22e-203
2 BOROUGH.1	0.422	0.00680	62.1	0.
3 BOROUGH.2	-0.414	0.00852	-48.5	0.
4 BOROUGH.4	-0.172	0.00586	-29.4	3.32e-188
5 BOROUGH.5	-0.299	0.00852	-35.0	1.95e-265
6 NEIGHBORHOODCIVIC.CENTER	0.578	0.0309	18.7	1.41e-77
7 NEIGHBORHOODTRIBECA	0.231	0.0232	9.95	2.59e-23
8 BLDG_CL_CATEGORY01.ONE.FAMILY.DWELLIN~	-0.0588	0.00897	-6.56	5.51e-11
9 BLDG_CL_CATEGORY07.RENTALS...WALKUP.A~	-0.0510	0.0207	-2.47	1.36e-2
10 BLDG_CL_CATEGORY08.RENTALS...ELEVATOR~	0.713	0.0573	12.4	1.61e-35
11 BLDG_CL_CATEGORY09.COOPS...WALKUP.APA~	0.0412	0.190	0.217	8.28e-1
12 BLDG_CL_CATEGORY10.COOPS...ELEVATOR.A~	0.129	0.0377	3.42	6.35e-4
13 BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.~	0.205	0.0107	19.2	6.28e-82
14 BLDG_CL_CATEGORY44.CONDO.PARKING	-1.76	0.0746	-23.6	3.62e-122
15 BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS~	-2.23	0.0641	-34.8	1.10e-262
16 GRSS_SF	0.258	0.00540	47.8	0.
17 TAX_CL_SOLD.1	-0.328	0.0128	-25.5	1.09e-142
18 TAX_CL_SOLD.2	-0.166	0.0162	-10.2	1.54e-24
19 BLDG_CL_SOLD1	0.445	0.0310	14.4	1.14e-46
20 BLDG_CL_SOLD6	-0.383	0.190	-2.02	4.38e-2
21 BLDG_CL_SOLD7	0.492	0.0412	11.9	9.03e-33
22 BLDG_CL_SOLD1	0.0937	0.0716	1.31	1.90e-1

23	BLDG_CL_SOLDD4	-0.362	0.0368	-9.82	9.39e- 23
24	BLDG_CL_SOLDRG	0.323	0.0809	3.99	6.67e- 5
25	SUM_UNIT	-0.0510	0.00572	-8.91	5.12e- 19

```
as.data.frame(varImp(model))%>%arrange(desc(Overall))
```

	Overall
BOROUGH.1	62.055
BOROUGH.2	48.529
GRSS_SF	47.841
BOROUGH.5	35.033
BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS.STORAGE	34.847
BOROUGH.4	29.395
TAX_CL_SOLD.1	25.521
BLDG_CL_CATEGORY44.CONDO.PARKING	23.575
BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS	19.210
NEIGHBORHOODCIVIC.CENTER	18.679
BLDG_CL_SOLDC1	14.361
BLDG_CL_CATEGORY08.RENTALS...ELEVATOR.APARTMENTS	12.449
BLDG_CL_SOLDC7	11.932
TAX_CL_SOLD.2	10.230
NEIGHBORHOODTRIBECA	9.953
BLDG_CL_SOLDD4	9.823
SUM_UNIT	8.914
BLDG_CL_CATEGORY01.ONE.FAMILY.DWELLINGS	6.558
BLDG_CL_SOLDRG	3.988
BLDG_CL_CATEGORY10.COOPS...ELEVATOR.APARTMENTS	3.416
BLDG_CL_CATEGORY07.RENTALS...WALKUP.APARTMENTS	2.467
BLDG_CL_SOLDC6	2.016
BLDG_CL_SOLDD1	1.309
BLDG_CL_CATEGORY09.COOPS...WALKUP.APARTMENTS	0.217

```
rmse_train_lm<-summary(model)$sigma

y_hat_lm<-predict.lm(model,newdata = test)
rmse_lm<-RMSE(y_hat_lm,test$SALE_LOG)

rmse_result <- data_frame(Method = "Linear Regression Predictor",RMSEtrain = rmse_train_lm,
                           RMSEtest = rmse_lm)
rmse_result%>%knitr::kable()
```

Method	RMSEtrain	RMSEtest
Linear Regression Predictor	0.464	0.475

The linear model generated an RMSE on the training set of .464 and on the test set of .475. The top 5 variables on this model in importance were BOROUGH.1 (Manhattan), BOROUGH.2 (Bronx), GRSS_SF, BOROUGH.5 (Staten Island), and BLDG_CL_CATEGORY47.CONDO.NON-BUSINESS.STORAGE.

rpart Model Moving onto a more complex algorithm, we next evaluate decision trees. Although the normalization/scaling that we completed on the dataset is not required for a decision tree, that step should

not have any negative impact on the *slower time to train* associated with this type of algorithm.

Typically with `rpart`, we have the opportunity to review the decision tree that produces the lowest RMSE results. In the case of our dataset, there are simply too many nodes and leaves to provide visual clarity, so we will instead provide the data on the points of distinction.

We have used a tuning option here on the complexity parameter (`cp`), the minimum benefit a split must add to the tree, and arrived at the `cp` of .002 as the ideal.

```
set.seed(1991,sample.kind="Rounding")
```

```
train_rpart <- train(SALE_LOG ~ ., method = "rpart",  
                    tuneGrid=data.frame(cp=seq(0,0.05,0.002)),data = train)  
train_rpart$bestTune
```

```
      cp  
2 0.002
```

```
train_rpart$finalModel
```

```
n= 46887
```

```
node), split, n, deviance, yval  
      * denotes terminal node
```

```
1) root 46887 17600.00  0.10300  
  2) BOROUGH.1< 0.5 35474 10800.00 -0.02490  
    4) GRSS_SF< 0.974 28438  6970.00 -0.11400  
      8) BLDG_CL_SOLDD4>=0.5 4949   617.00 -0.39700  
        16) BOROUGH.2>=0.5 765    107.00 -0.60800 *  
        17) BOROUGH.2< 0.5 4184   470.00 -0.35800 *  
    9) BLDG_CL_SOLDD4< 0.5 23489  5880.00 -0.05450  
      18) BLDG_CL_CATEGORY44.CONDO.PARKING>=0.5 217    163.00 -1.41000 *  
      19) BLDG_CL_CATEGORY44.CONDO.PARKING< 0.5 23272  5310.00 -0.04190  
        38) BLDG_CL_CATEGORY09.COOPS...WALKUP.APARTMENTS>=0.5 1387   225.00 -0.39800  
          76) BOROUGH.4>=0.5 928    40.90 -0.50400 *  
          77) BOROUGH.4< 0.5 459   153.00 -0.18400  
            154) BOROUGH.2>=0.5 40    10.80 -1.15000 *  
            155) BOROUGH.2< 0.5 419   101.00 -0.09110 *  
    39) BLDG_CL_CATEGORY09.COOPS...WALKUP.APARTMENTS< 0.5 21885  4900.00 -0.01930  
      78) BOROUGH.2>=0.5 1953   539.00 -0.27800  
        156) TAX_CL_SOLD.2>=0.5 262    80.90 -0.64800 *  
        157) TAX_CL_SOLD.2< 0.5 1691   417.00 -0.22100 *  
    79) BOROUGH.2< 0.5 19932  4220.00  0.00601  
      158) BOROUGH.5>=0.5 3908   670.00 -0.18500 *  
      159) BOROUGH.5< 0.5 16024  3370.00  0.05270  
        318) BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS< 0.5 12729  2800.00  0.00715  
          636) BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS.STORAGE>=0.5 25    4.19 -2.07000 *  
          637) BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS.STORAGE< 0.5 12704  2690.00  0.01120  
            1274) GRSS_SF< 0.67 8886  1910.00 -0.03150  
              2548) BOROUGH.4>=0.5 5465  1020.00 -0.08530 *  
              2549) BOROUGH.4< 0.5 3421   853.00  0.05450 *  
                1275) GRSS_SF>=0.67 3818   721.00  0.11100 *  
    319) BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS>=0.5 3295   443.00  0.22900
```

```

        638) BOROUGH.4>=0.5 1366    116.00  0.06970 *
        639) BOROUGH.4< 0.5 1929    269.00  0.34100 *
5) GRSS_SF>=0.974 7036    2650.00  0.33500
10) GRSS_SF< 1.91 6151    1600.00  0.23100
    20) BOROUGH.2>=0.5 1036     222.00 -0.06290 *
    21) BOROUGH.2< 0.5 5115    1270.00  0.29000
        42) GRSS_SF< 1.34 3565     760.00  0.23100 *
        43) GRSS_SF>=1.34 1550     467.00  0.42700 *
11) GRSS_SF>=1.91 885     516.00  1.06000
    22) GRSS_SF< 2.98 537     204.00  0.81900 *
    23) GRSS_SF>=2.98 348     231.00  1.44000
        46) BLDG_CL_CATEGORY10.COOPS...ELEVATOR.APARTMENTS>=0.5 20      9.39 -0.40200 *
        47) BLDG_CL_CATEGORY10.COOPS...ELEVATOR.APARTMENTS< 0.5 328    150.00  1.55000 *
3) BOROUGH.1>=0.5 11413    4430.00  0.50100
6) GRSS_SF< 1.4 10808    3430.00  0.44800
12) BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS< 0.5 6199    1830.00  0.27600
24) BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS.STORAGE>=0.5 29      7.62 -1.57000 *
25) BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS.STORAGE< 0.5 6170    1720.00  0.28500
50) BLDG_CL_SOLD6>=0.5 594     120.00 -0.00338 *
51) BLDG_CL_SOLD6< 0.5 5576    1550.00  0.31600
    102) BLDG_CL_SOLD4>=0.5 4118     905.00  0.26300 *
    103) BLDG_CL_SOLD4< 0.5 1458     598.00  0.46300 *
13) BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS>=0.5 4609    1170.00  0.68000
26) NEIGHBORHOODCIVIC.CENTER< 0.5 4415    1070.00  0.65300 *
27) NEIGHBORHOODCIVIC.CENTER>=0.5 194      28.10  1.28000 *
7) GRSS_SF>=1.4 605     429.00  1.44000
14) BLDG_CL_CATEGORY10.COOPS...ELEVATOR.APARTMENTS>=0.5 7      4.95 -0.68400 *
15) BLDG_CL_CATEGORY10.COOPS...ELEVATOR.APARTMENTS< 0.5 598    392.00  1.47000
30) GRSS_SF< 4.79 556     310.00  1.39000
60) GRSS_SF>=4.76 16      2.50 -1.21000 *
61) GRSS_SF< 4.76 540     196.00  1.46000 *
31) GRSS_SF>=4.79 42      29.50  2.55000 *

```

```
varImp(train_rpart)
```

rpart variable importance

only 20 most important variables shown (out of 24)

	Overall
GRSS_SF	100.00
BLDG_CL_SOLD4	52.38
BLDG_CL_CATEGORY10.COOPS...ELEVATOR.APARTMENTS	51.23
SUM_UNIT	39.68
BOROUGH.2	37.17
BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS	35.14
BOROUGH.4	26.80
BLDG_CL_CATEGORY08.RENTALS...ELEVATOR.APARTMENTS	20.56
TAX_CL_SOLD.1	19.25
TAX_CL_SOLD.2	16.29
BOROUGH.5	13.10
BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS.STORAGE	11.40
BLDG_CL_SOLD6	10.71
BLDG_CL_CATEGORY44.CONDO.PARKING	10.56

```
BLDG_CL_CATEGORY09.COOPS...WALKUP.APARTMENTS    10.55
BOROUGH.1                                         8.99
BLDG_CL_SOLDRG                                   6.00
NEIGHBORHOODTRIBECA                             4.26
NEIGHBORHOODCIVIC.CENTER                        4.17
BLDG_CL_SOLDL1                                   2.62
```

```
rmse_train_rpart<-train_rpart$results$RMSE[which.min(train_rpart$results$RMSE)]

y_hat_rpart <- predict(train_rpart, test, type = "raw")
rmse_rpart<-RMSE(y_hat_rpart,test$SALE_LOG)

rmse_result <- bind_rows(rmse_result,
                        data_frame(Method="rpart cp = .002",
                                   RMSEtrain = rmse_train_rpart, RMSEtest = rmse_rpart))

rmse_result %>% knitr::kable()
```

Method	RMSEtrain	RMSEtest
Linear Regression Predictor	0.464	0.475
rpart cp = .002	0.464	0.461

The rpart model generated an RMSE on the training set of .464, and on the test set of .461. The top 5 variables in this model, in order of importance, were GRSS_SF, BLDG_CL_SOLDL4, BLDG_CL_CATEGORY10.COOPS...ELEVATOR APARTMENTS, SUM_UNIT, and BOROUGH.2 (Bronx).

randomForest Model Random Forest as an algorithm gives us the ability to have R process many alternative combinations of variables in order to determine the combination that delivers the lowest RMSE. This algorithm is more robust than the single decision tree rpart algorithm we just fit.

For this particular model, and in the interest of processing time, we chose 64 trees (*informed by How Many Trees in a Random Forest? Authors Thais Mayumi Oshiro, Pedro Santoro Perez, José Augusto Baranauskas*), and tuned the mtry, the number of variables randomly sampled in each tree, which produced a bestTune mtry metric of 11. The higher mtry reduces the correlation between trees in the forest, thus improving prediction accuracy, but the tuning process and the *slower time to fit* associated with random Forest make this segment the longest to run of the three models we fit. Although mtry values of 1-14 were tested, we have abbreviated the code to 9-12 in this model to minimize run time.

```
set.seed(14,sample.kind = "Rounding")

fit_rf <- train(SALE_LOG ~ ., method = "rf",
               tuneGrid = data.frame(mtry = 9:12),ntree=64, data = train)
fit_rf$bestTune

  mtry
3    11

fit_rf
```

Random Forest

46887 samples
24 predictor

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 46887, 46887, 46887, 46887, 46887, 46887, ...

Resampling results across tuning parameters:

mtry	RMSE	Rsquared	MAE
9	0.451	0.459	0.289
10	0.451	0.460	0.289
11	0.450	0.460	0.289
12	0.451	0.459	0.289

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was mtry = 11.

```
varImp(fit_rf)
```

rf variable importance

only 20 most important variables shown (out of 24)

	Overall
GRSS_SF	100.00
BOROUGH.1	73.54
BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS	21.80
SUM_UNIT	19.57
BLDG_CL_SOLDD4	15.97
BOROUGH.2	13.75
BLDG_CL_CATEGORY10.COOPS...ELEVATOR.APARTMENTS	11.69
BLDG_CL_CATEGORY44.CONDO.PARKING	11.68
BLDG_CL_CATEGORY47.CONDO.NON.BUSINESS.STORAGE	7.98
BOROUGH.4	7.26
BOROUGH.5	6.54
BLDG_CL_CATEGORY09.COOPS...WALKUP.APARTMENTS	5.57
BLDG_CL_SOLDC6	5.33
TAX_CL_SOLD.1	4.63
TAX_CL_SOLD.2	3.96
BLDG_CL_SOLDRG	3.46
NEIGHBORHOODCIVIC.CENTER	2.33
BLDG_CL_CATEGORY08.RENTALS...ELEVATOR.APARTMENTS	1.97
BLDG_CL_CATEGORY07.RENTALS...WALKUP.APARTMENTS	1.63
NEIGHBORHOODTRIBECA	1.12

```
rmse_train_rf<-fit_rf$results$RMSE[which.min(fit_rf$results$RMSE)]
y_hat_rf<-predict(fit_rf, test, type = "raw")
rmse_rf<-RMSE(y_hat_rf,test$SALE_LOG)
rmse_result <- bind_rows(rmse_result,
                          data_frame(Method="randomForest ntree=64, mtry=11",
                                      RMSEtrain = rmse_train_rf,
```

```

RMSEtest = rmse_rf))
rmse_result %>% knitr::kable()

```

Method	RMSEtrain	RMSEtest
Linear Regression Predictor	0.464	0.475
rpart cp = .002	0.464	0.461
randomForest ntree=64, mtry=11	0.450	0.452

The random Forest model generated an RMSE of .450 on the training set, and .452 on the test set.

Reviewing the variable importance, we note that gross square footage of the property (*GRSS_SF*) was a variable present in all of the decision trees in the ensemble, which is interesting considering that it was not, as discussed earlier, the variable with the highest correlation to SALE LOG. Rounding out the top 5 variables, in order of importance: BOROUGH.1 (Manhattan), BLDG_CL_CATEGORY13.CONDOS...ELEVATOR.APARTMENTS, SUM_UNIT, and BLDG_CL_SOLDD4.

Final Model and Results

Final Model - The machine learning model trained that resulted in the *lowest RMSE of .452* was the randomForest algorithm with mtry=11 as noted above. Of the 24 predictors that were used in the model, only the top 8 were included in more than 10% of the trees in the ensemble, and as noted above, gross square footage (**GRSS_SF**) was present in each of the models.

Conclusion

The RMSE results of all three models fit ranged from .475 to .452, supporting our original premise that we could build a machine learning algorithm that could predict Sale Price with relative accuracy based on an evaluation of other variables in the dataset such as Borough, Square Footage, Building Class, and Building Class Category. The random Forest algorithm specifically should be the one for future application given the RMSE results.

These models should now be used to evaluate future 12 month period of New York City Property Sales. While it would be my assumption that there may be macro level trends that alter the effectiveness of this model for future periods, it would be useful to understand whether the key predictors of this dataset remain the key predictors of future time periods, even if the degree of correlation varies. Conceivably, this model, after evaluation of additional years of data, could be a key tool for real estate developers to use in predicting the return on investment over time on new property development.

Limitations of the model, next steps

Limitations of the model

The RMSE results here are tied to this model's fit to the specific subset of data that was partitioned in this process into the **train** set(training set) and **test** set (test set), and might not be replicated as favorably should the same model be used on future data partitions of the same **nyc** dataset.

Other limitations of the model are that it is trained on a relatively small set of 24 variables, and that the period of data is only 12 months. It serves well on this dataset in isolation. Additionally, as is true of any model, the quality of the model is contingent on the consistency and accuracy of the data therein.

Future Work

Interesting future work on this dataset could include:

- revising the data partitioning to allow for k fold cross validation (60% training, 20% cross validation, 20% testing hold out)
- putting the original dataset without log transformation or standardization through the model fitting process to understand the benefit of that original process in either processing time or RMSE output
- constructing an ensemble of these machine learning algorithms
- converting the dataset to a matrix focused on the combination of borough, block, and lot (commonly called a BBL), which forms a unique key for property in New York City. Block and Lot were removed from the original dataset due to low direct correlation with Sale Price, but building a separate model against BBL could include stratification by groupings of block number, and would be another intriguing analysis on this same dataset